
Bivalent Representations for Meta-Reinforcement Learning

Gonalo Guiomar

Abstract

asdasda

1 Introduction

Basic points to put here

- Exploration of high-dimensional state-spaces is costly
- Having a compact representation is necessary to efficiently transverse environments
- Reward signals are not symmetric in frequency: positive rewards are sparse and localized and negative rewards are diffused
- The brain processes positive and negative reward information through separate pathways
- How can we combine this into a more efficient RL algorithm?
- We combine graph representation learning techniques with reward sensitivity functions in order to speed up learning

2 Background

3 Splitting Representations

3.1 Successor Features

3.2 Reward sensitivities

3.3 Combining multiple representations

4 Experiments

5 Conclusion