
Contextual Learning with Reward Sensitive Representations

Gonçalo Guiomar

Tiago Costa

Margarida Sousa

Abstract

...

1 Introduction

Basic points to put here:

- Exploration of high-dimensional state-spaces is costly
- Having a compact representation is necessary to efficiently transverse environments
- Reward signals in nature are not symmetric in frequency of occurrence: positive rewards are sparse and localized and negative rewards are diffused
- These signals are also contextual; animals seem to bootstrap policies of new environments on what has been learned in other environments
- Besides, the brain processes positive and negative reward information through separate pathways
- How can we combine this into a more efficient RL algorithm?
- We'll attempt to define a reward and context sensitive algorithm that learns policies over a set of environments
- We combine graph representation learning techniques with reward sensitivity functions in order to speed up learning across different MDPs

2 Background

3 Feature Representations over Graphs

4 Contextual Learning

We begin by defining a state value function that takes into account the sensitivity to certain reward ranges, as in [REF-Tano] we write

$$V^h(s_t) = \mathbf{E}^\pi \left[\sum_{\tau} \gamma^\tau f^h(r_{t+\tau}) \right] \quad (1)$$

Although feature vectors allow for the compression of state relevant information in a way that allows for very high-dimensional state-spaces to be explored, we lose a bit of a notion of context for each state; to which extent is the state that the agent is in at time t related contextually with states that have a similar representation?

To explore this idea we define a novel value function as dependent on the dot product of the state embeddings in a surrounding ball of size ϵ by weighting the surrounding weights with a function $\Gamma^\epsilon(s_t)$

$$\hat{V}^h(s_t) = \sum_i \Gamma^\epsilon(s_t)_i \langle \phi_i, \mathbf{w}^h \rangle \quad (2)$$

with $\langle \phi_i, \mathbf{w}^h \rangle$ the inner product between the weights \mathbf{w} and the embeddings of each state ϕ_i

$$\Gamma^\epsilon(s_t)_i = \begin{cases} 0 & \text{if } \|\phi(s_t) - \phi_i\| > \epsilon \\ \frac{1}{1 + \beta_V \|\phi(s_t) - \phi_i\|} & \text{otherwise} \end{cases} \quad (3)$$

and the loss function for each RPE function is given by

$$\mathcal{L}^h = \|V^h(s_t) - \hat{V}^h(s_t)\|^2 \quad (4)$$

which if we take the following expansion

$$V^h(s_t) = \mathbf{E}^\pi \left[f^h(r_t) + \gamma \sum_{\tau} \gamma^\tau f^h(r_{t+1+\tau}) \right] = f^h(r_t) + V^h(s_{t+1}) \quad (5)$$

allows us to write the following objective function

$$\mathcal{L}^h = \|f^h(r_t) + \gamma V^h(s_{t+1}) - \hat{V}^h(s_t)\|^2. \quad (6)$$

which implies new definition for the RPE

$$\delta_t^h = f^h(r_t) + \gamma V^h(s_{t+1}) - \hat{V}^h(s_t). \quad (7)$$

We can learn the weights \mathbf{w} by the usual gradient descent rule

$$\frac{\partial \mathcal{L}^h}{\partial w_j^h} = 2\delta_t^h \frac{\partial \delta_t^h}{\partial w_j^h} \quad (8)$$

where in calculating the latter partial derivative we need to take into account that only the approximated value function can be differentiated in respect to \mathbf{w} so

$$\frac{\partial \delta_t^h}{\partial w_j^h} = -\frac{\partial \hat{V}^h(s_t)}{\partial w_j^h} = -\sum_i \Gamma^\epsilon(s_t)_i \phi_i \quad (9)$$

which will give us a final update rule for each weight w_j

$$w_j \leftarrow w_j - \alpha^h 2\delta_t^h \sum_i \Gamma^\epsilon(s_t)_i \phi_i \quad (10)$$

4.1 Policy

As a first approach we can try to model the policy with an action preference formalism that takes the RPEs of each channel h and feeds it into a corresponding preference

$$A^h(s, a) \leftarrow A^h(s, a) + \alpha \delta_t^h \quad (11)$$

where we sum all channels

$$A^T(s, a) = \sum_h \theta_h A^h(s, a) \quad (12)$$

and define our policy through a softmax over the total

$$\pi(a|s, \theta) = \frac{e^{\sum_h \theta_h A^h(s, a)}}{\sum_a e^{\sum_h \theta_h A^h(s, a)}} \quad (13)$$

5 Experiments

6 Conclusion